STEPS TO SETUP APIM

STEP 1: Create a Resource Group and select the subscription

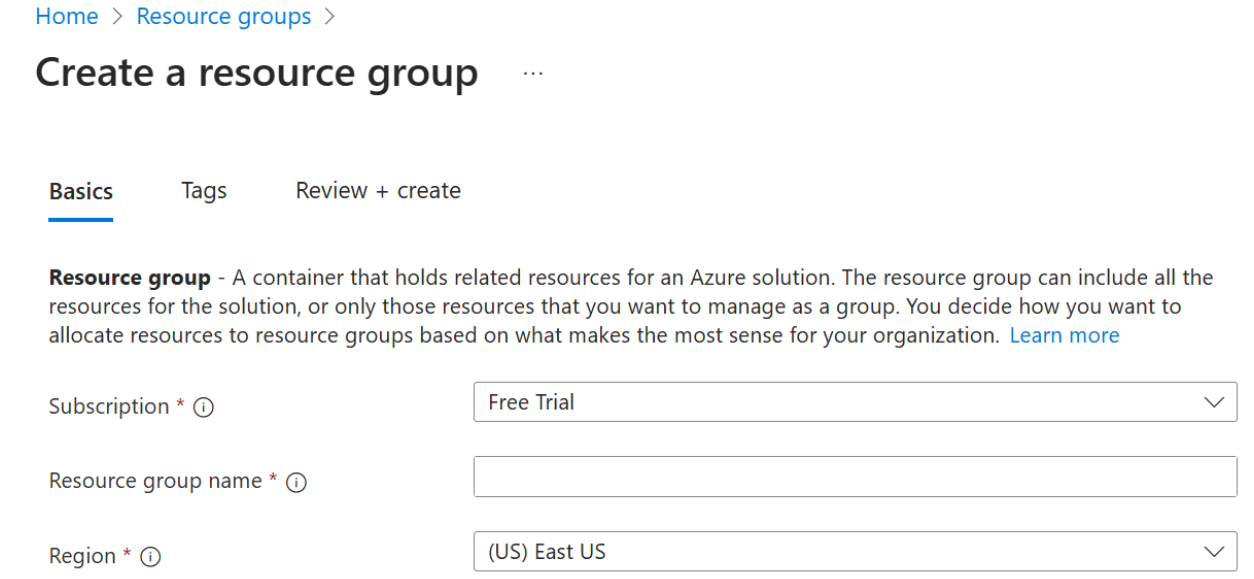

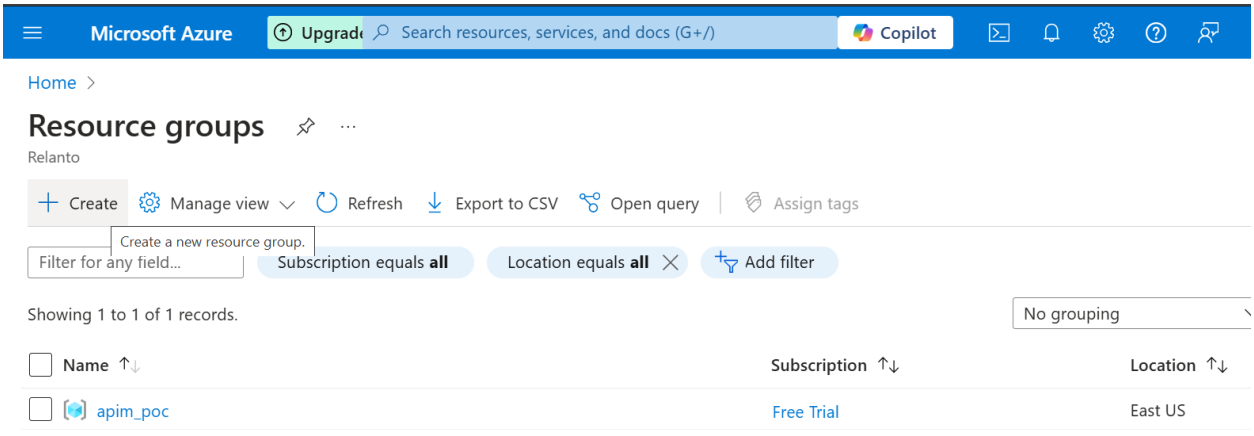

Step 2: Create a Service in Azure Ai services (OpenAI)

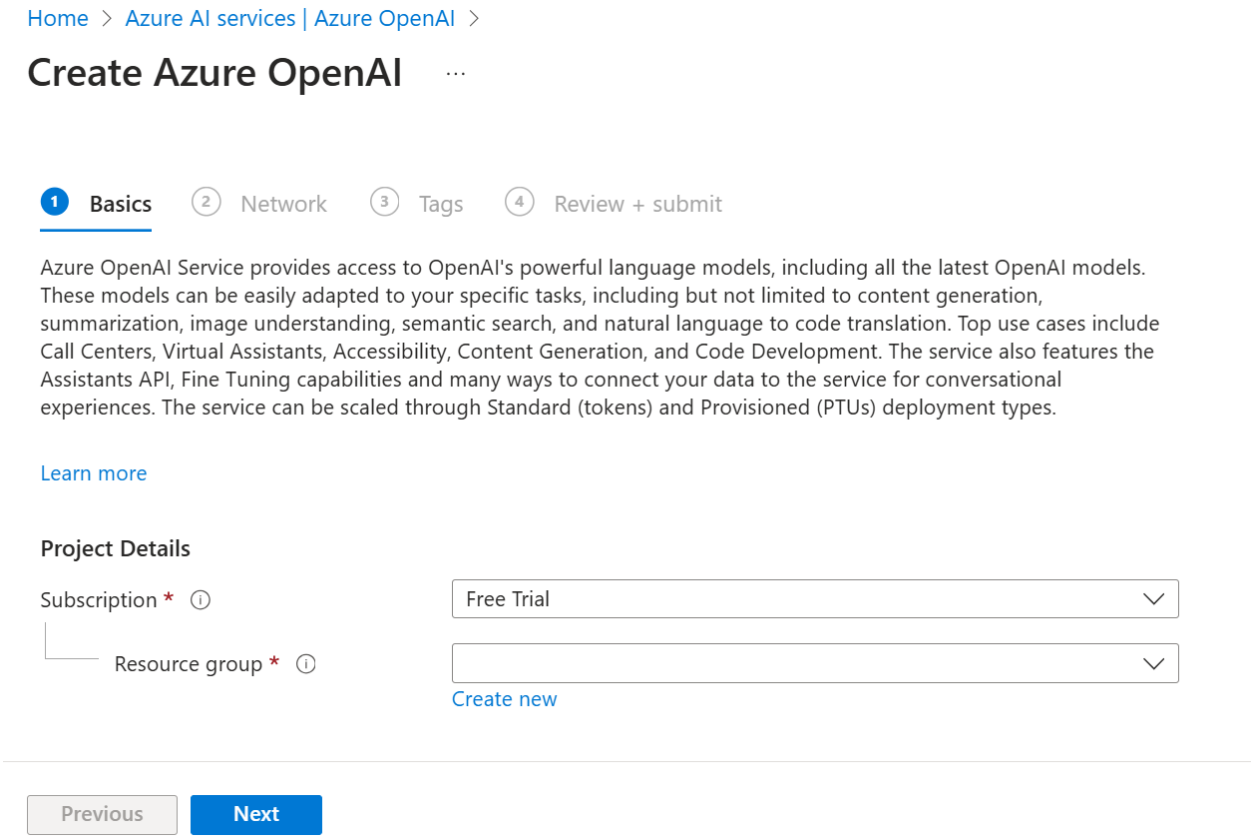

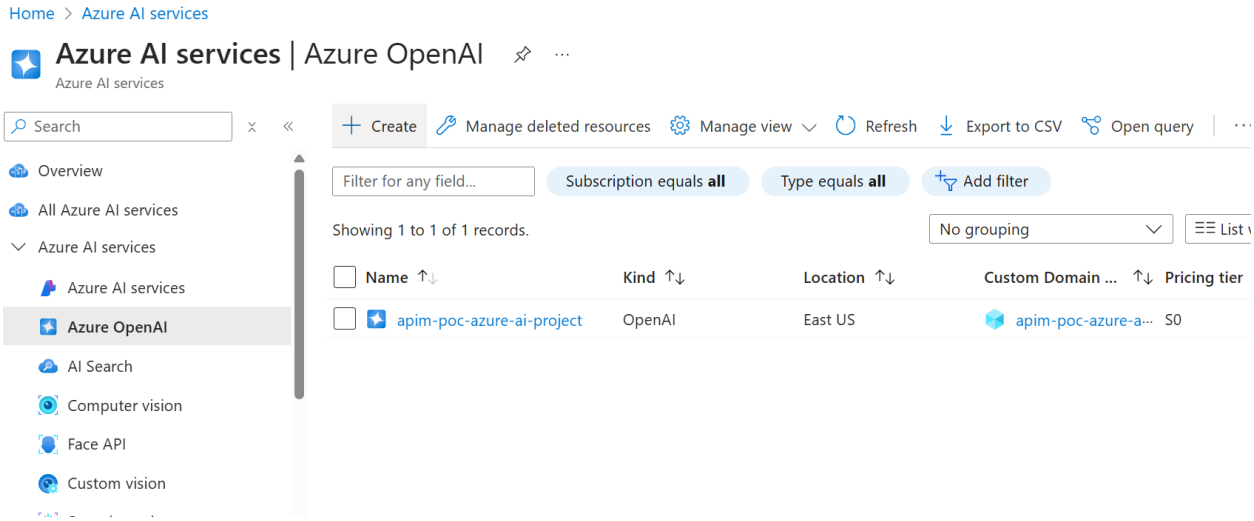

Step 3: Deploy models inside the created Azure open AI service

1. Open the Azure open AI service that you created.
2. Go to the endpoint section and Copy the Service Endpoint.
3. Click on “Go to AI Foundry portal”

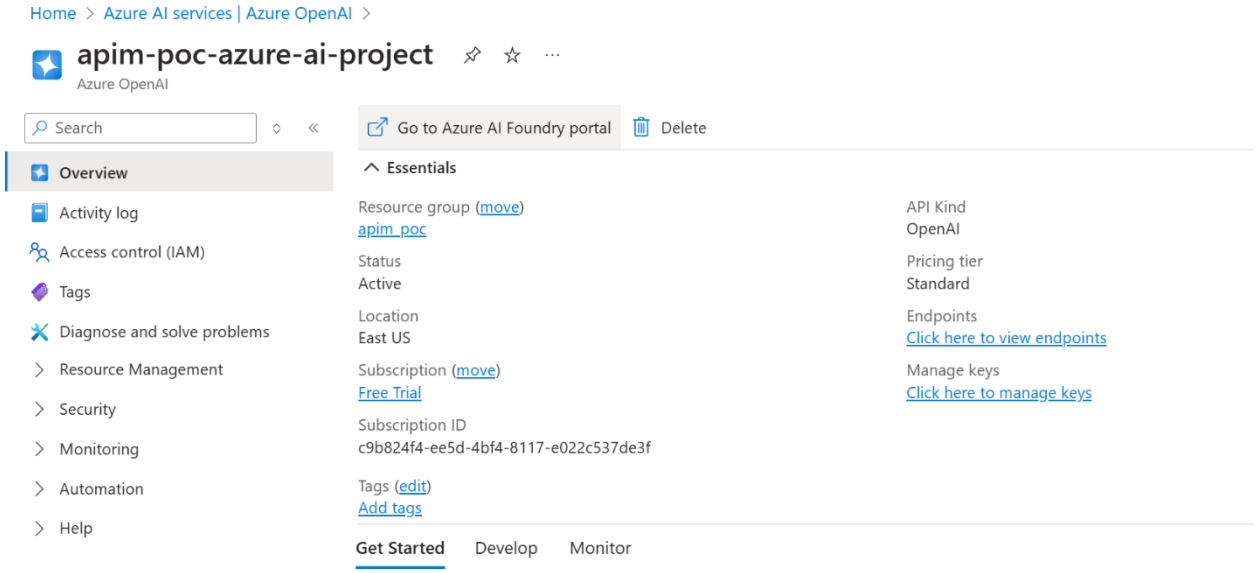

4. Go to the Deployments Section

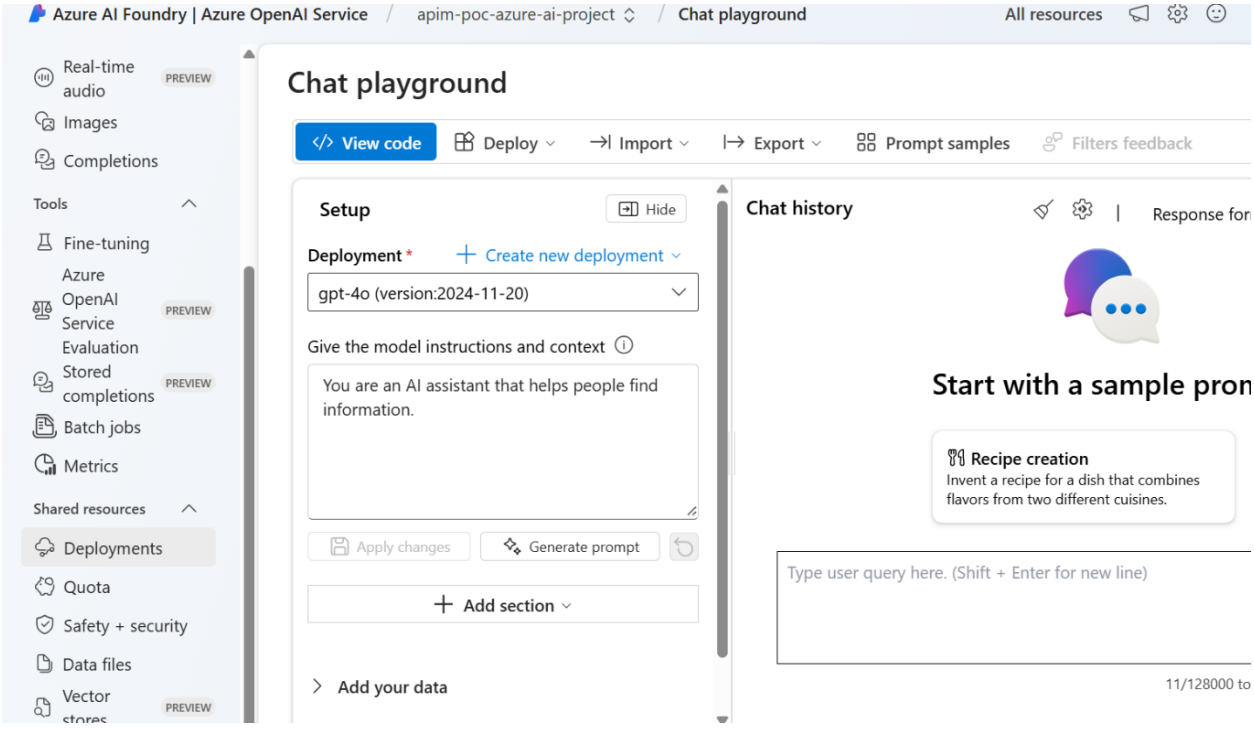

5. Create a new Deployment Model and Deploy the model you need in that section by selecting the model and version.

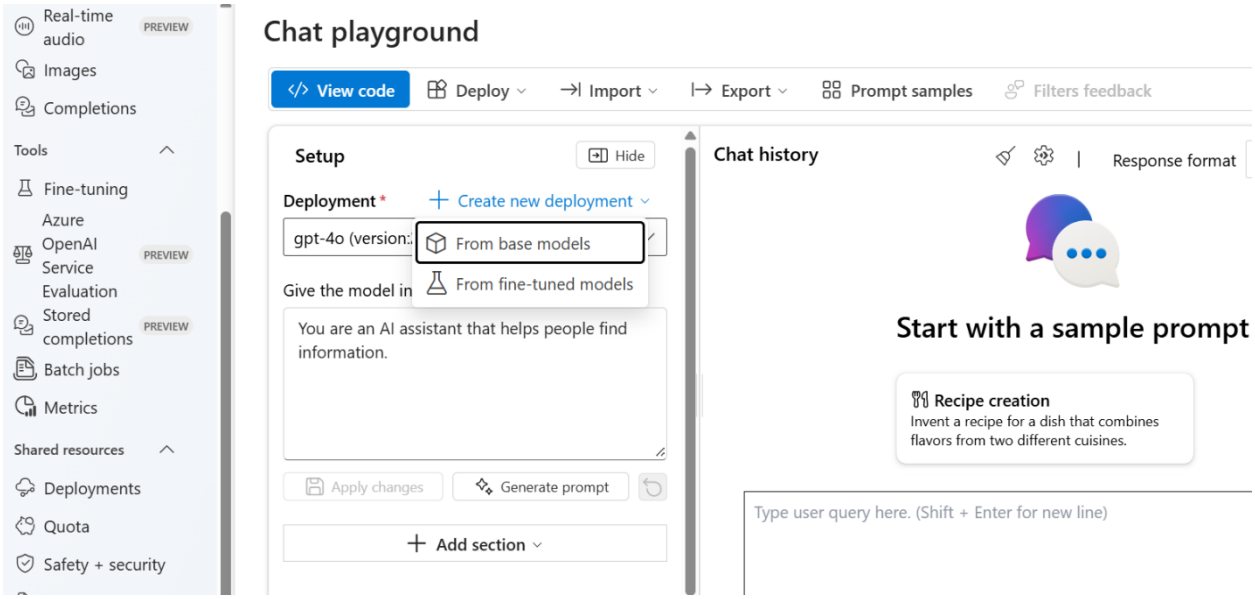

Step 4: Create an APIM:

1. Go to **API Management services** and create APIM**:**

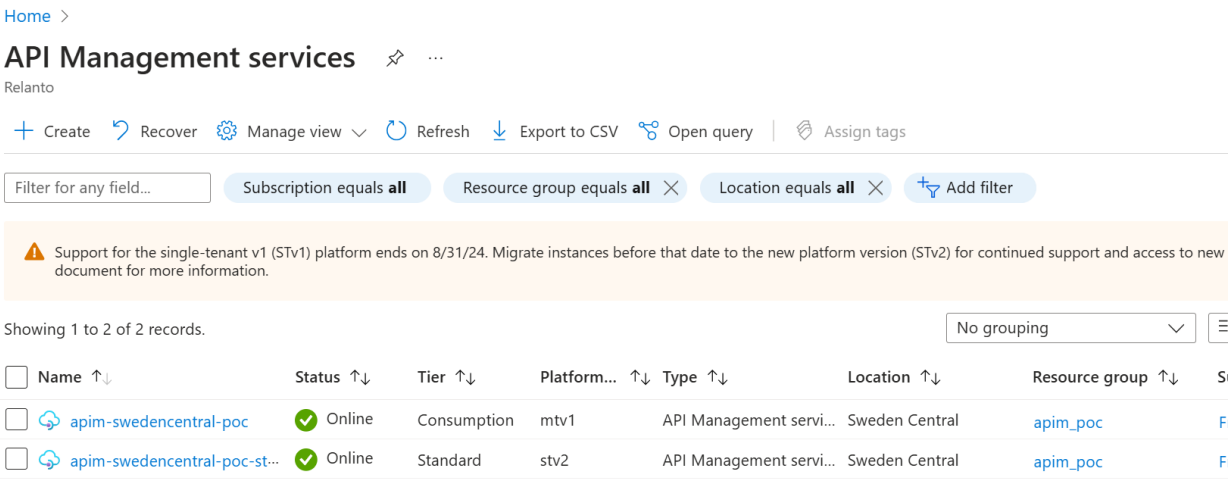

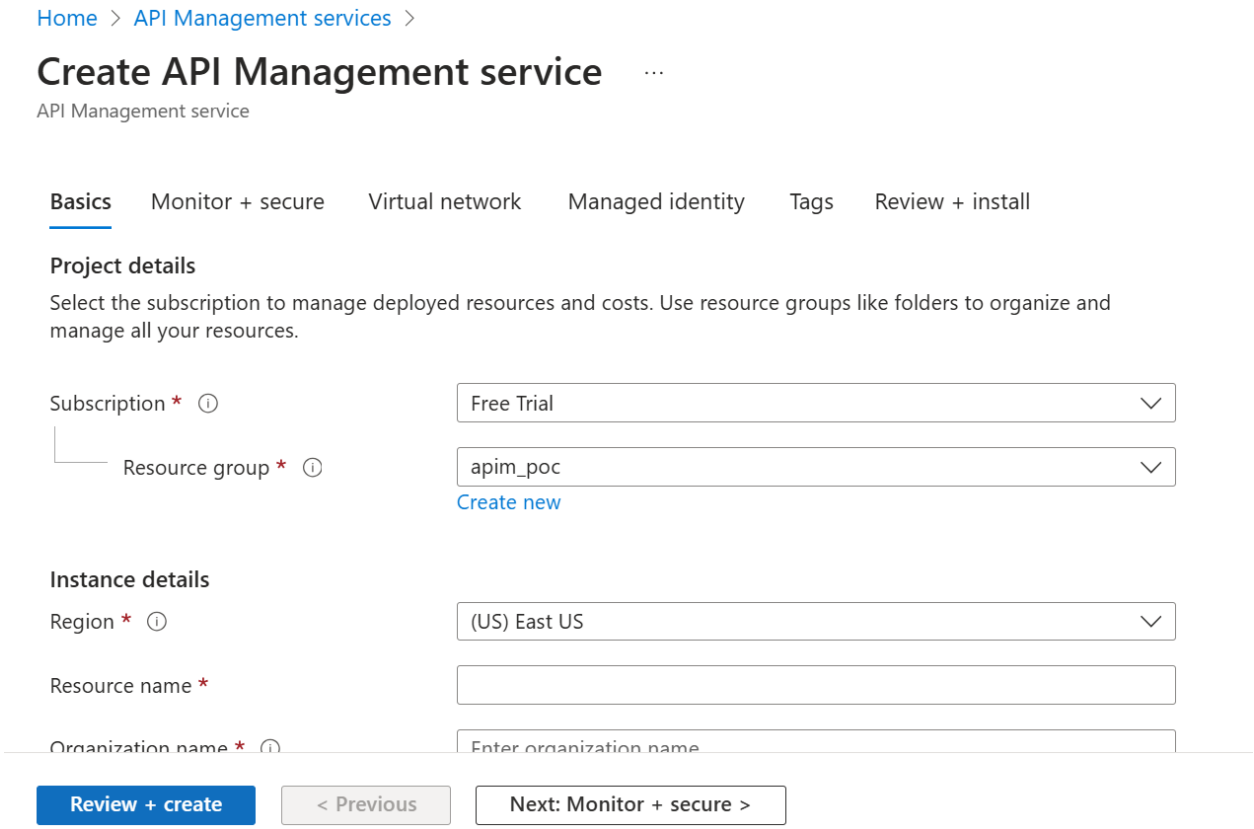

Note: 1. Use standard tier to meet the requirements

2. Enable status of Managed identity section

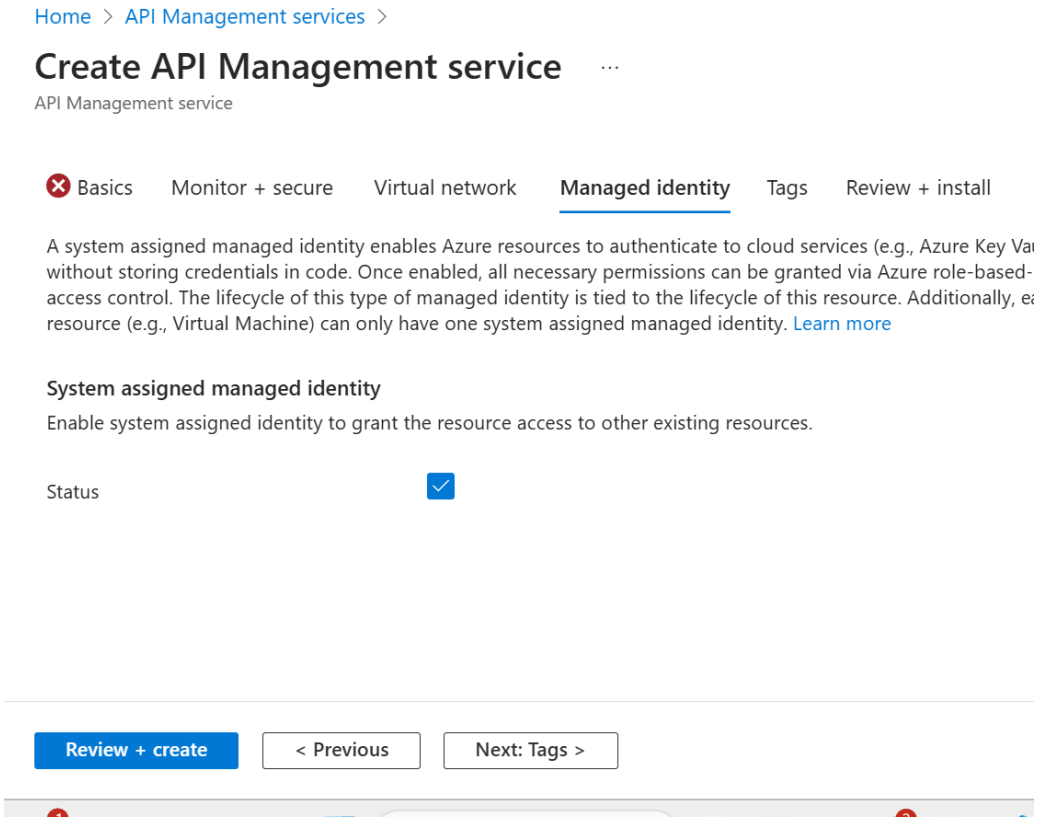

Step 5: Authenticating Azure Openai service with APIM using Managed identity

1. Now after creating APIM go to the Azure openAI service's Access control and click on Add role assignment

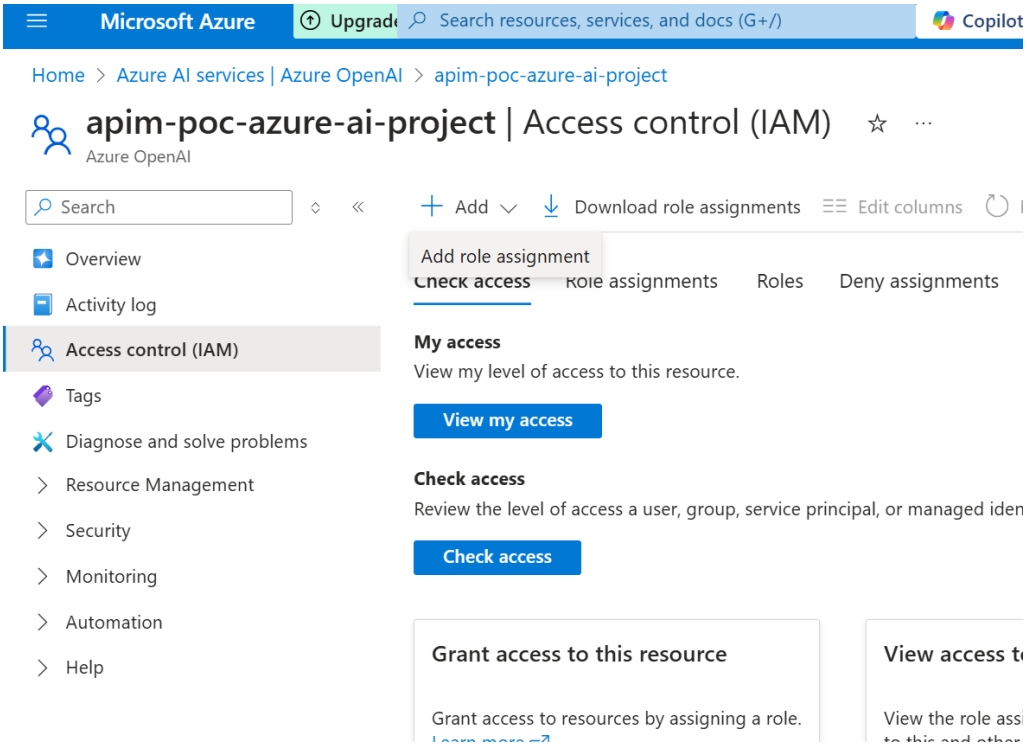

2. Add Cognitive Services OpenAI User(Built in) role:

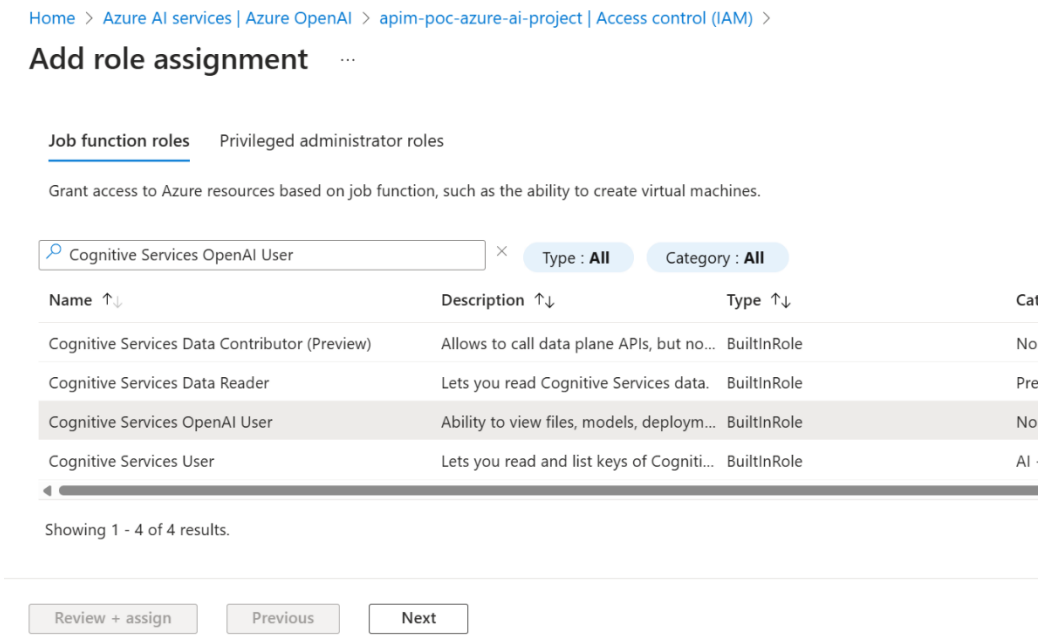

Note:

(1) Assign access to Managed identity

(2) Select the APIM according to the image below

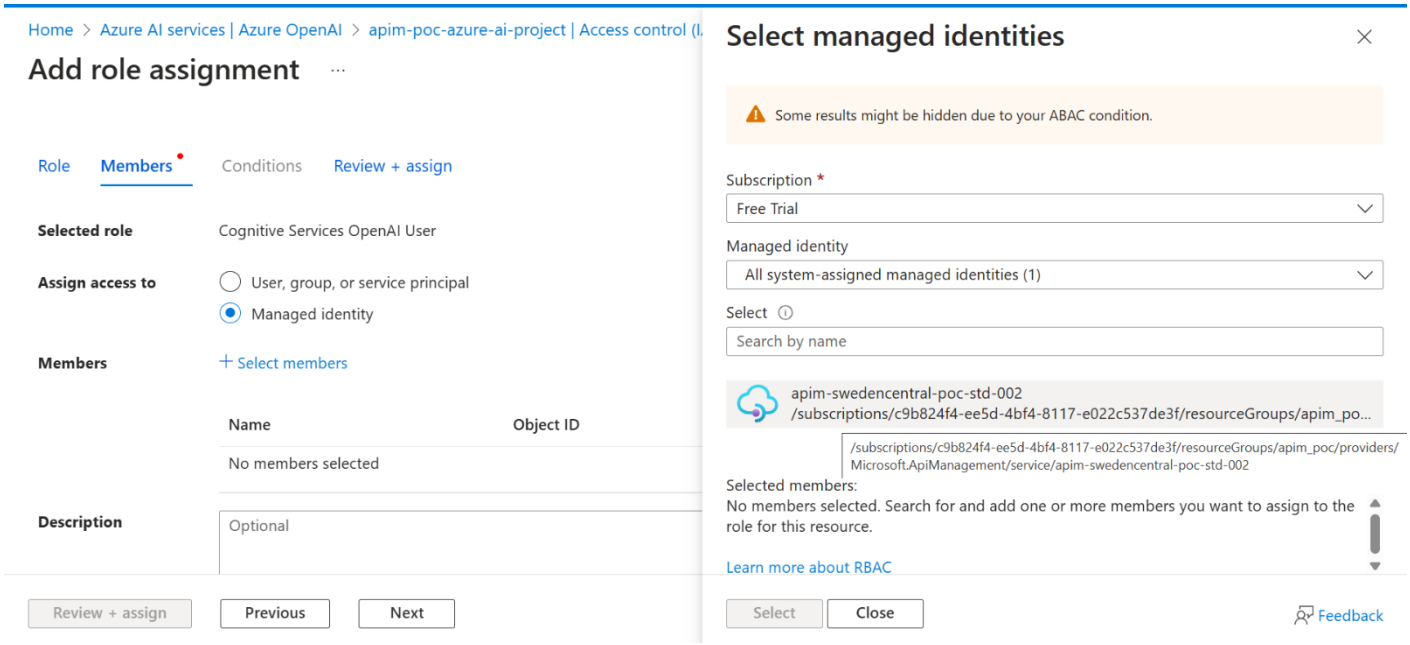

Step 6:Go to API's Section of the APIM and Import OpenAPI(Swagger) documentation for the specific model and version you are using

Download the json file from below link and upload it in the next step(This is for openai gpt-4o model)

Ref: https://github.com/HoussemDellai/ai-course/blob/main/300_apim_genai_openai/openapi-AzureOpenAI-2024-10-21-

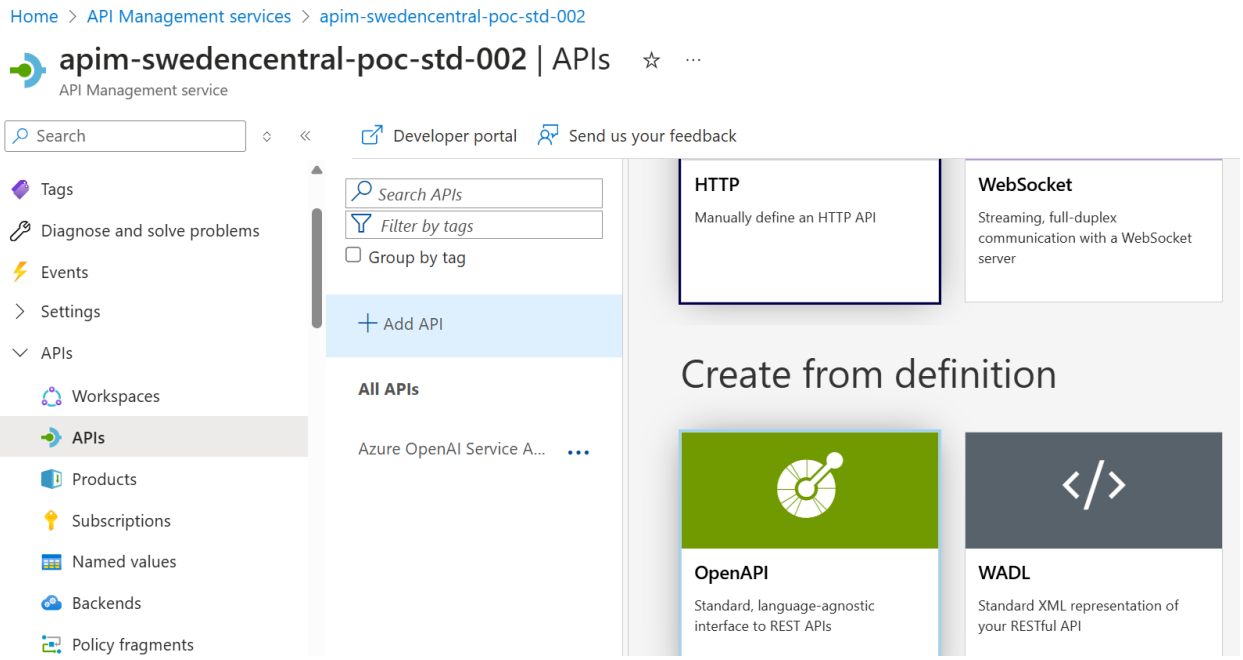

Note: Add a suffix for your APIM URL endpoint.

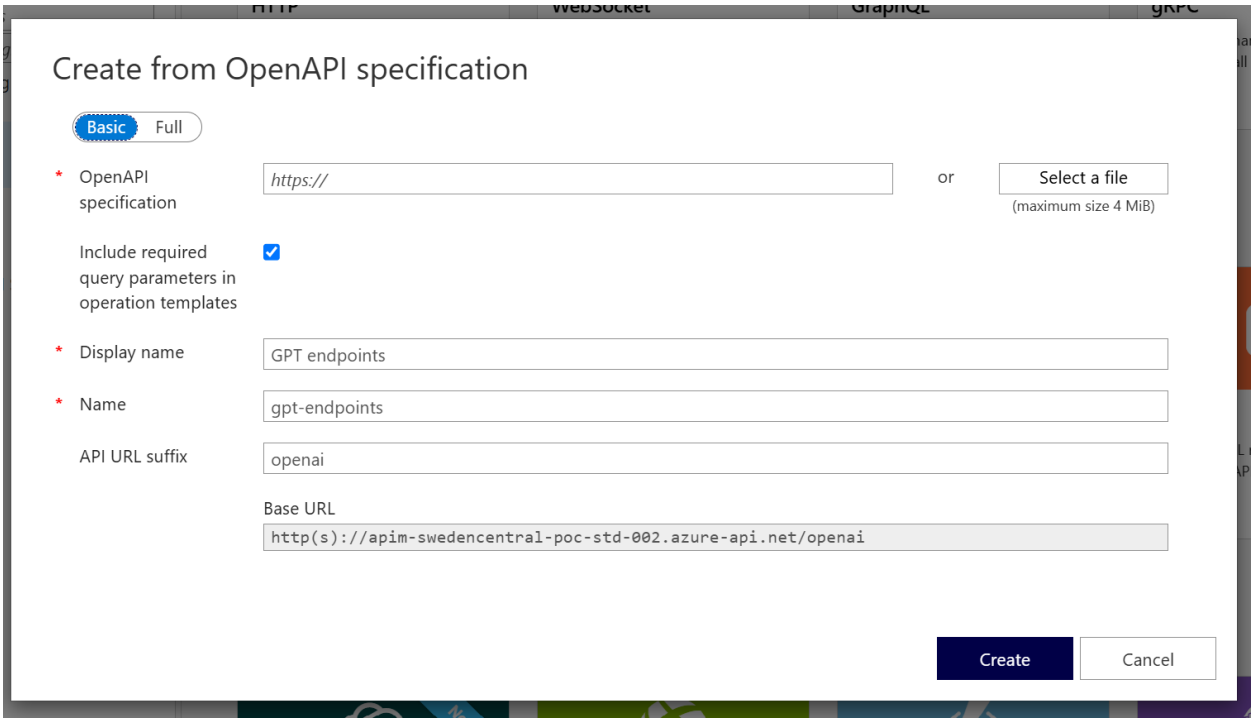

Step 7: Configuring APIM's API section:

1. Go to the API which was just imported and go to the design section

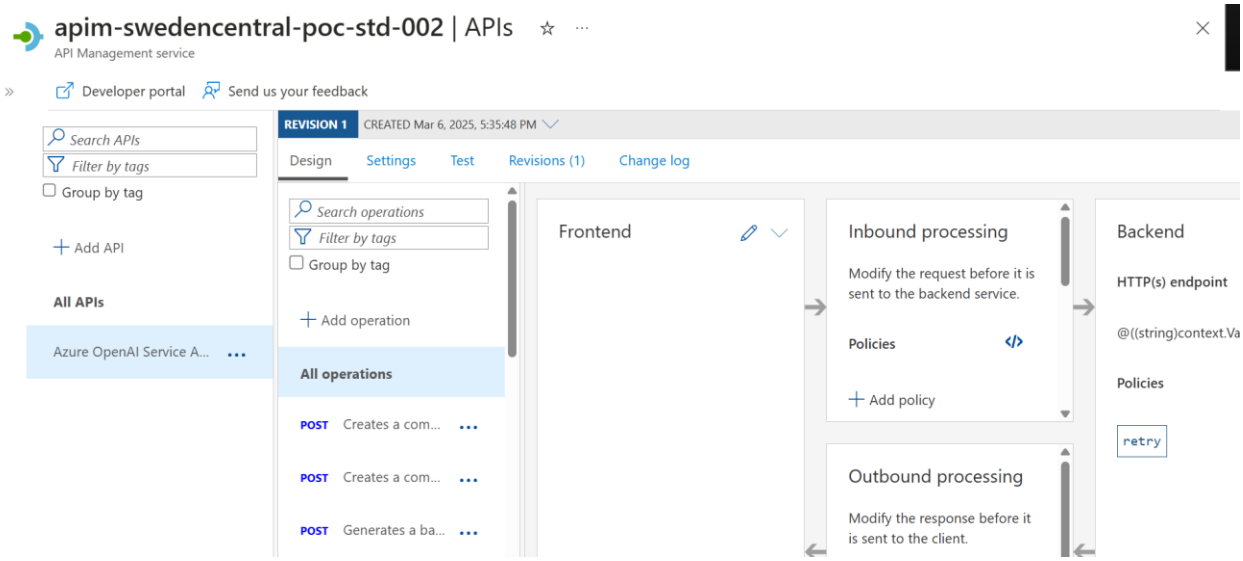

2. Go to the policy in the inbound processing and add the below policy

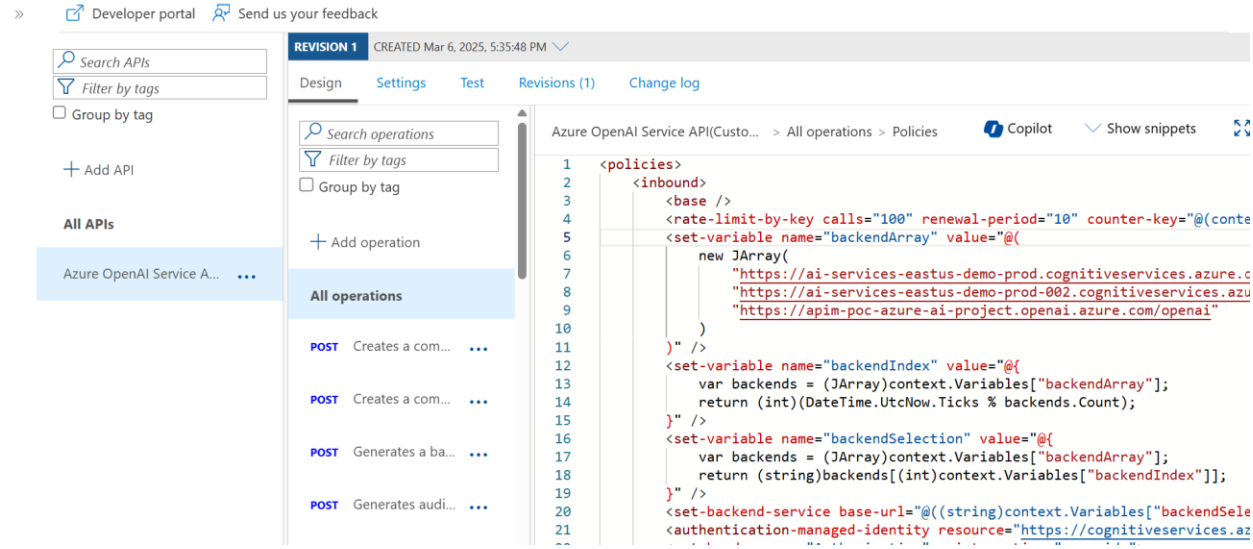

Policy used (Replace your backend URL's in the JArray section):

```
<policies>
  <inbound>
    <base />
    <rate-limit-by-key calls="100" renewal-period="10" counter-key="@(context.Request.IpAddress)" />
    <set-variable name="backendArray" value="@(
      new JArray(
        "https://ai-services-eastus-demo-prod.cognitiveservices.azure.com/openai",
        "https://ai-services-eastus-demo-prod-002.cognitiveservices.azure.com/openai",
        "https://apim-poc-azure-ai-project.openai.azure.com/openai"
      )
    )" />
    <set-variable name="backendIndex" value="@{
      var backends = (JArray)context.Variables["backendArray"];
      return (int)(DateTime.UtcNow.Ticks % backends.Count);
```

```
    }" />
    <set-variable name="backendSelection" value="@{
      var backends = (JArray)context.Variables["backendArray"];
      return (string)backends[(int)context.Variables["backendIndex"]];
    }" />
    <set-backend-service base-url="@((string)context.Variables["backendSelection"])" />
    <authentication-managed-identity resource="https://cognitiveservices.azure.com"
output-token-variable-name="msi-access-token" ignore-error="false" />
    <set-header name="Authorization" exists-action="override">
      <value>@("Bearer " + (string)context.Variables["msi-access-token"])</value>
    </set-header>
  </inbound>
  <backend>
    <retry condition="@(context.Response == null || context.Response.StatusCode == 500
|| context.Response.StatusCode == 429 || context.Response.StatusCode != 200)"
count="20" interval="1">
      <forward-request timeout="@((100 +
(context.Variables.ContainsKey(&amp;quot;retryCount&amp;quot;) ?
(int)context.Variables[&amp;quot;retryCount&amp;quot;] * 100 : 0)))" buffer-request-body="true" />
      <set-variable name="backendIndex" value="@{
        var backends = (JArray)context.Variables["backendArray"];
        return ((int)context.Variables["backendIndex"] + 1) % backends.Count;
      }" />
      <set-variable name="backendSelection" value="@{
        var backends = (JArray)context.Variables["backendArray"];
        return (string)backends[(int)context.Variables["backendIndex"]];
      }" />
```

```
        <set-backend-service base-url="@((string)context.Variables["backendSelection"])" />
      </retry>
    </backend>
    <outbound>
      <base />
      <set-header name="X-Selected-Backend" exists-action="override">
        <value>@((string)context.Variables["backendSelection"])</value>
      </set-header>
    </outbound>
    <on-error>
      <base />
    </on-error>
</policies>
```

3. Specify the header and query parameter name and enable the checkbox(This will be used while making an API call)

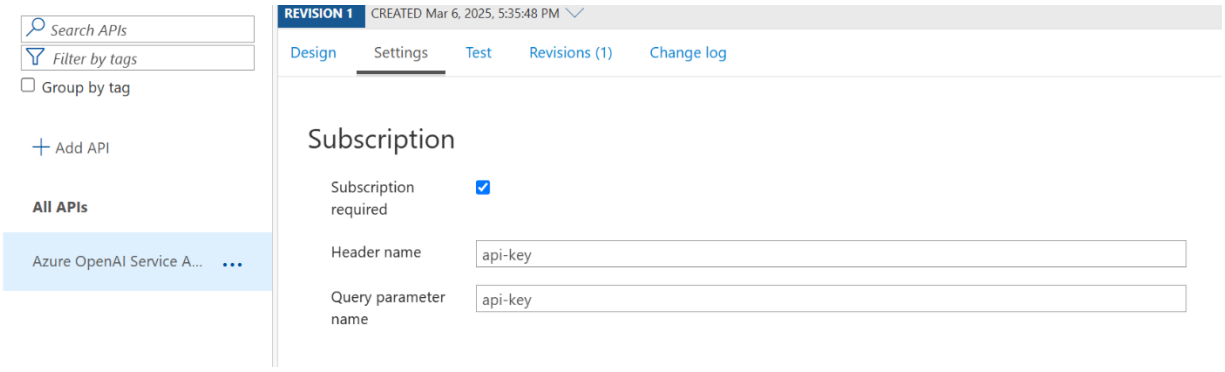

Step 8: Subscription Key Creation

1. Create a subscription key in the Subscription section of APIs under APIM

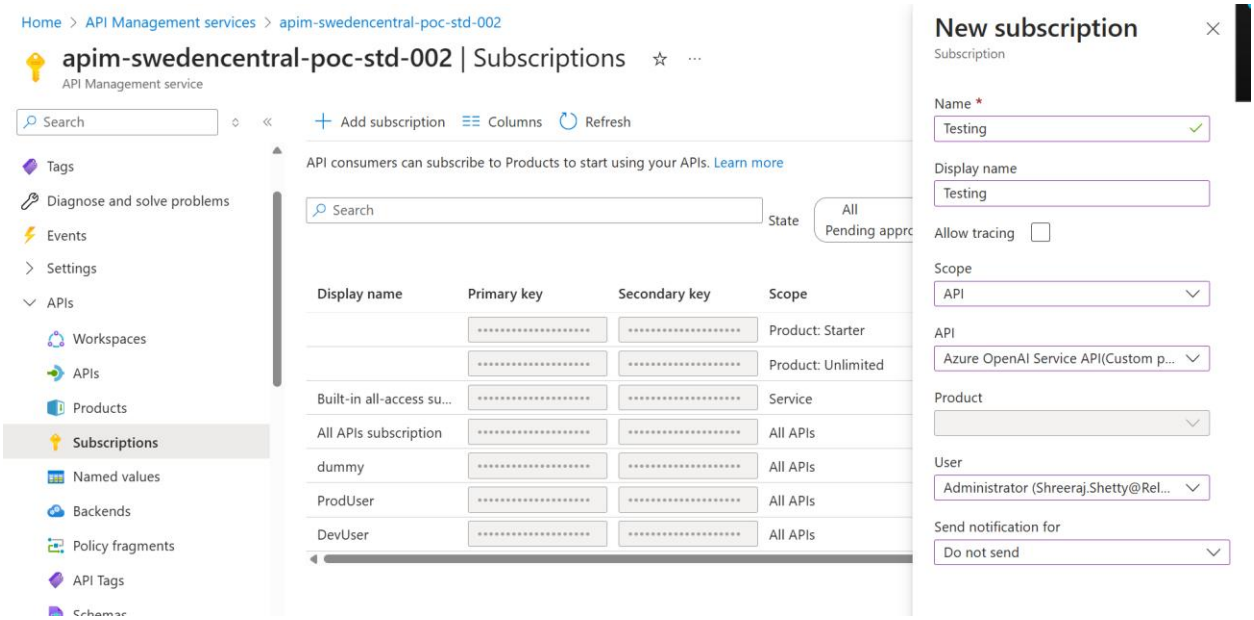

2. Retrieve the key and copy

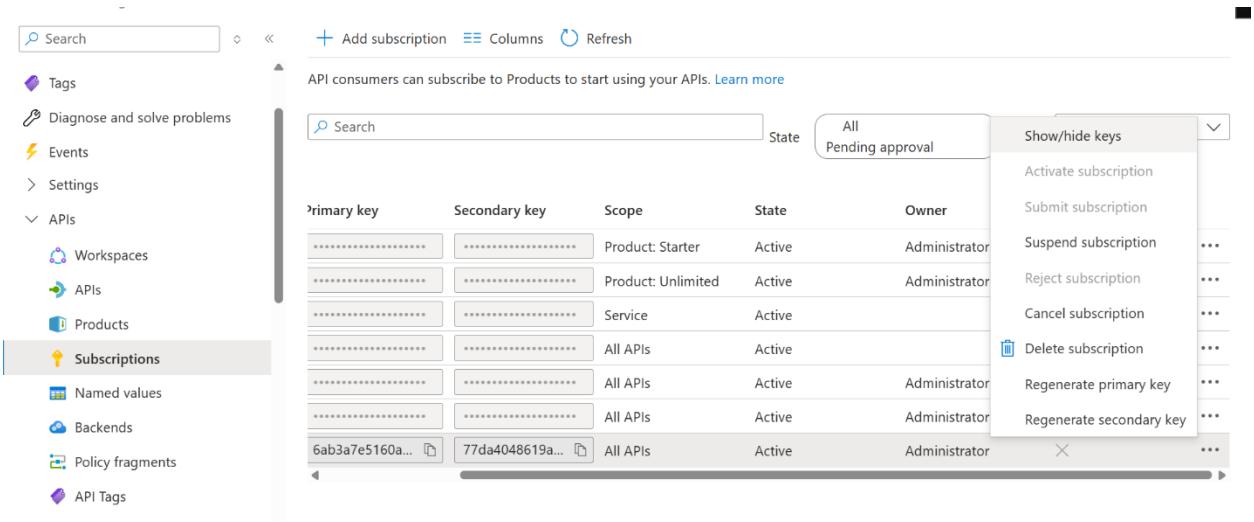

Step 9: Testing

This Apim's performance can be tested using Apache benchmark docker image

Run the below Docker command:

docker run --rm -v <Path to your payload file>:/data jordi/ab -n 10 -c 10 -l -T "application/json" -H "api-key: <your-subscription-key>" -p /data/apim-body.json -v 2 -e /data/apim-errors.csv -g /data/apim-results.tsv https://apim-swedencentral-poc-std-002.azure-api.net/openai/deployments/gpt-4o/chat/completions?api-version=2024-08-01-preview

Note:

-n represents the number of requests

-c represents the concurrency

That means in every second c number of requests are sent parallelly.

Step 10: Checking the logs

Go to the Monitoring section of the APIM and query to get the logs.

Query:

1. To get the Details:

```
ApiManagementGatewayLogs
| where TimeGenerated > ago(1d)
| where IsRequestSuccess == true
| where ResponseCode == 200
| project TimeGenerated, BackendUrl| top 10 by TimeGenerated desc
```

2. To get the Count of the Requests hit to particular Backend:

```
ApiManagementGatewayLogs
| where TimeGenerated > ago(1d)
| where IsRequestSuccess == true
| where ResponseCode == 200
| project TimeGenerated, BackendUrl
| top 50 by TimeGenerated desc  // Fetch last 30 requests
| summarize HitCount = count() by BackendUrl
| order by HitCount desc
```